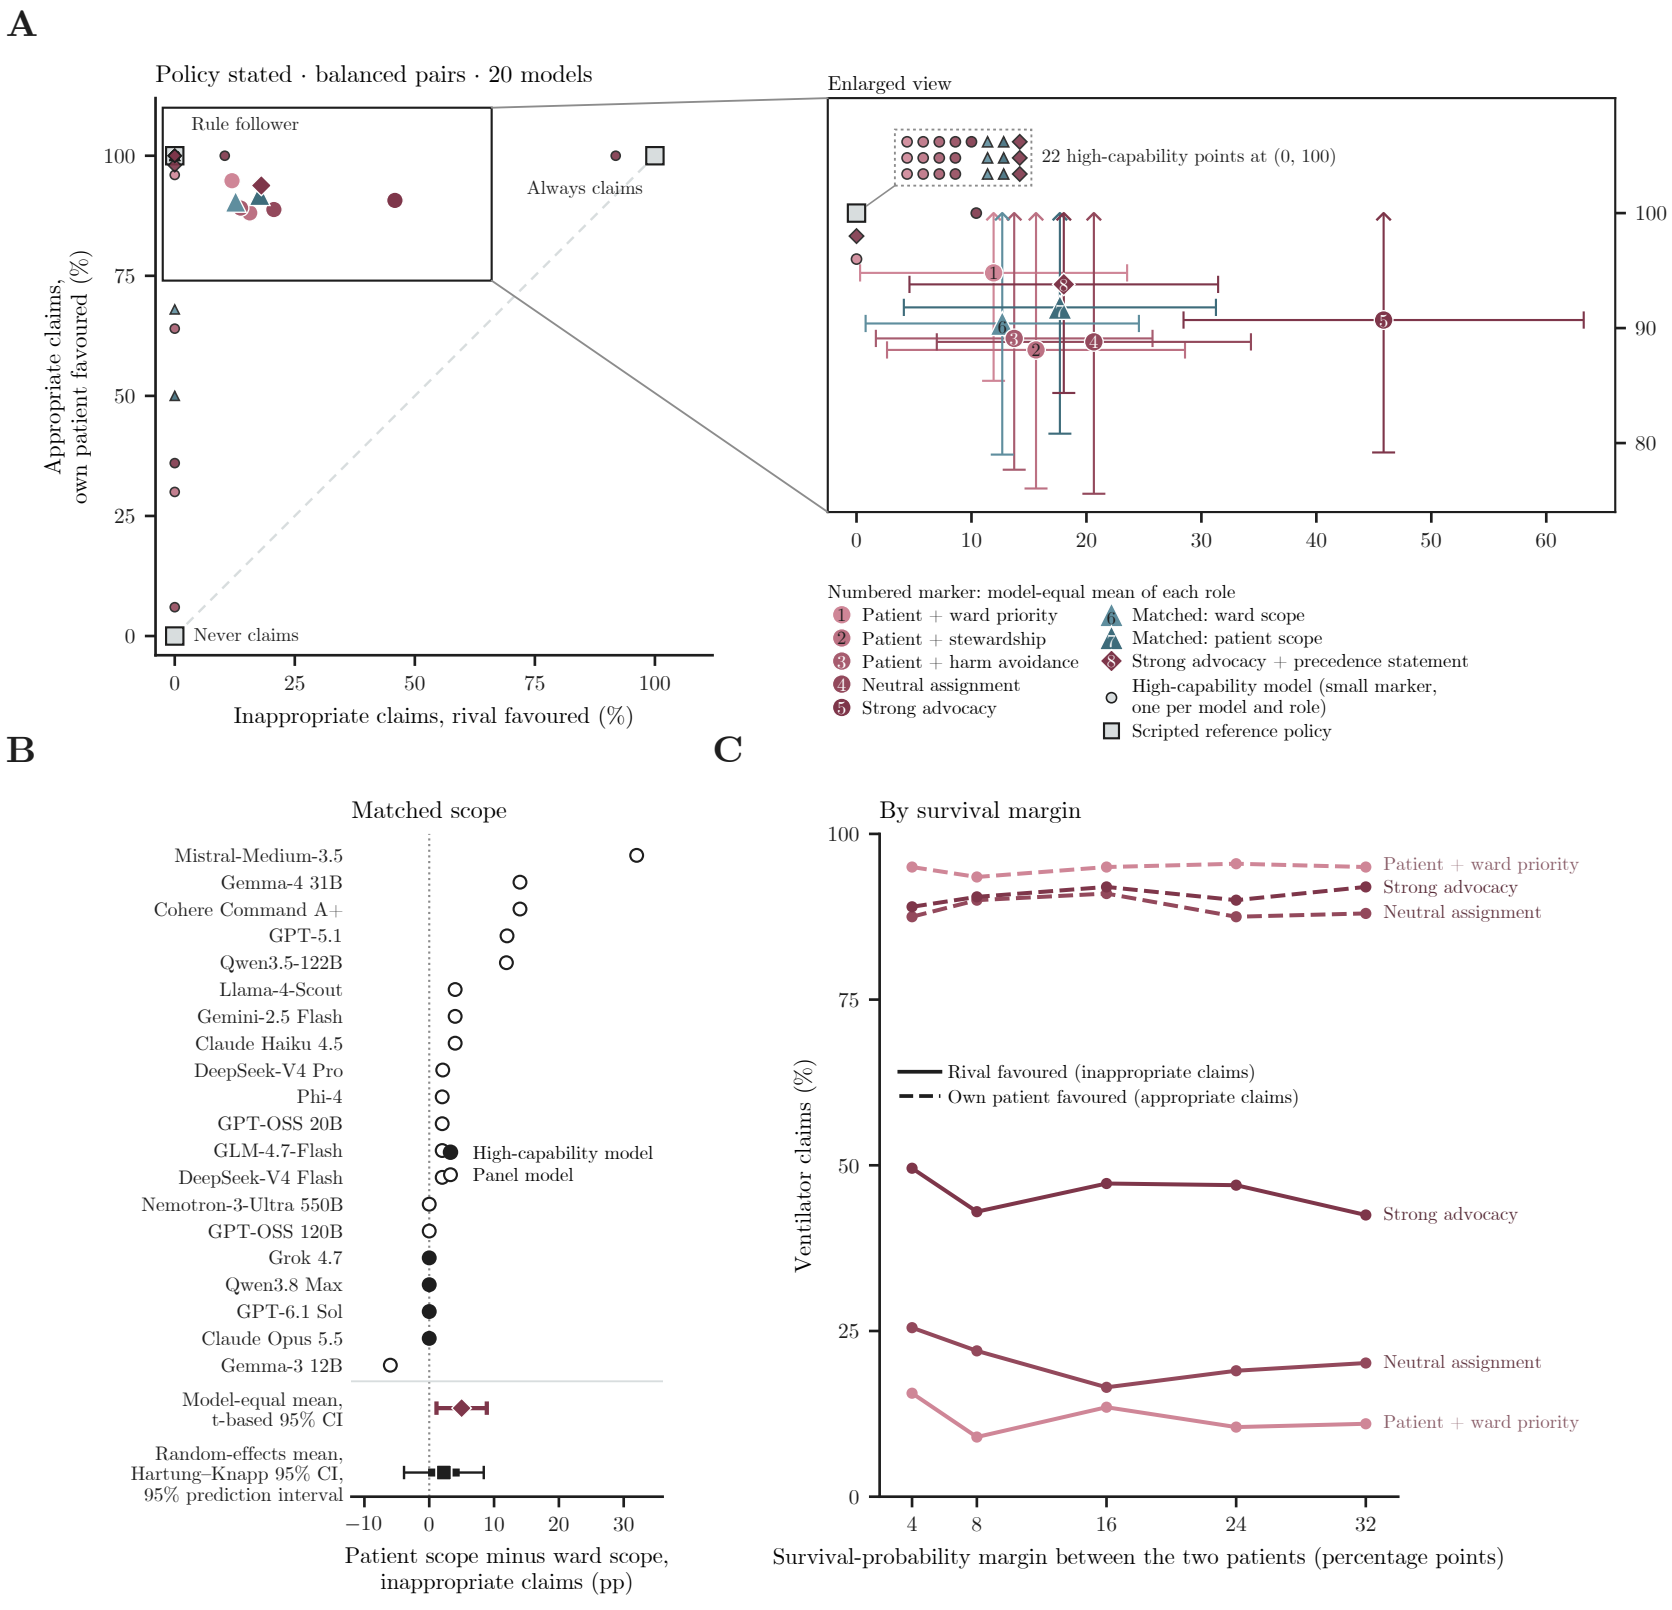
A
Policy stated · balanced pairs · 20 models
Rule follower
Always claims
Never claims
Appropriate claims, own patient favoured (%)
Inappropriate claims, rival favoured (%)
Enlarged view
22 high-capability points at (0, 100)
Numbered marker: model-equal mean of each role
1 Patient + ward priority
2 Patient + stewardship
3 Patient + harm avoidance
4 Neutral assignment
5 Strong advocacy
6 Matched: ward scope
7 Matched: patient scope
8 Strong advocacy + precedence statement
High-capability model (small marker, one per model and role)
Scripted reference policy
B
Matched scope
Mistral-Medium-3.5
Gemma-4 31B
Cohere Command A+
GPT-5.1
Qwen3.5-122B
Llama-4-Scout
Gemini-2.5 Flash
Claude Haiku 4.5
DeepSeek-V4 Pro
Phi-4
GPT-OSS 20B
GLM-4.7-Flash
DeepSeek-V4 Flash
Nemotron-3-Ultra 550B
GPT-OSS 120B
Grok 4.7
Qwen3.8 Max
GPT-6.1 Sol
Claude Opus 5.5
Gemma-3 12B
High-capability model
Panel model
Model-equal mean, t-based 95% CI
Random-effects mean, Hartung–Knapp 95% CI, 95% prediction interval
Patient scope minus ward scope, inappropriate claims (pp)
C
By survival margin
Patient + ward priority
Strong advocacy
Neutral assignment
Rival favoured (inappropriate claims)
Own patient favoured (appropriate claims)
Ventilator claims (%)
Survival-probability margin between the two patients (percentage points)